

Step 1: Setup Google Cloud Project

✓ In Google Cloud Console:

1. Go to **Cloud Console** → **Project Selector** → **New Project**

Name it, for example:

perl
CopyEdit
my-rag-agent-project

- 2.
3. Click **Create**.

Step 2: Enable Required APIs

In **Cloud Console** → **APIs & Services** → **Enabled APIs**:

- Enable these APIs:
 - Vertex AI API
 - Discovery Engine API.
 - Cloud Storage API
 - Identity and Access Management API

If you don't see "Agentspace", you may have to request allowlist from Google sales, as it is enterprise-featured.

Step 3: Create a Cloud Storage Bucket

✓ In Cloud Storage:

1. Go to **Storage** → **Buckets** → **Create**

Name it for example:

perl
CopyEdit
my-agentspace-bucket

- 2.
3. Choose a region (say, **us-central1**)
4. Click **Create**.

✓ Upload your PDF:

- In that bucket, click **Upload File** → pick your **employee_policy.pdf**.

The screenshot shows the Google Cloud Storage console for the 'rag-agent-project'. The left sidebar contains navigation links for Cloud Storage, Overview, Buckets, Monitoring, Settings, Storage Intelligence, Insights datasets, and Configuration. The main panel displays the 'Bucket details' for 'my-agentspace-bucket'. The 'Objects' tab is active, showing a folder browser view with the bucket name 'my-agentspace-bucket'. Below this, there are buttons for 'Create folder', 'Upload', 'Transfer data', and 'Other services'. A table lists the objects in the bucket, filtered by name prefix only. The table has columns for Name, Size, Type, Created, and Storage class. The objects listed are:

Name	Size	Type	Created	Storage class
An_Effective_Retrieval_Method_to...	270.6 KB	application/pdf	Jul 6, 2025, 2:48:23 PM	Standard
NeurIPS-2020-retrieval-augmented...	946.1 KB	application/pdf	Jul 6, 2025, 2:48:25 PM	Standard
borgeaud22a.pdf	986.9 KB	application/pdf	Jul 6, 2025, 2:48:26 PM	Standard
paper1.pdf	6.3 MB	application/pdf	Jul 6, 2025, 2:48:40 PM	Standard
paper2.pdf	199.2 KB	application/pdf	Jul 6, 2025, 2:48:30 PM	Standard
paper3.pdf	11 MB	application/pdf	Jul 6, 2025, 2:48:43 PM	Standard
yulius-h-2020-forex-market-screen...	4.2 MB	application/pdf	Jul 6, 2025, 2:48:50 PM	Standard

Go to Discovery Engine

In Google Cloud Console, go to the left side menu:

nginx

CopyEdit

Vertex AI → Discovery Engine (or Vertex AI → Agent Builder / AgentSpace, depending on your console name)

1. If you do not see Discovery Engine, you might need to enable *Discovery Engine API* in APIs & Services.
2. In Discovery Engine, choose:
Data Stores → Create Data Store
Since you want to do RAG over PDF documents (unstructured knowledge base), you should select:

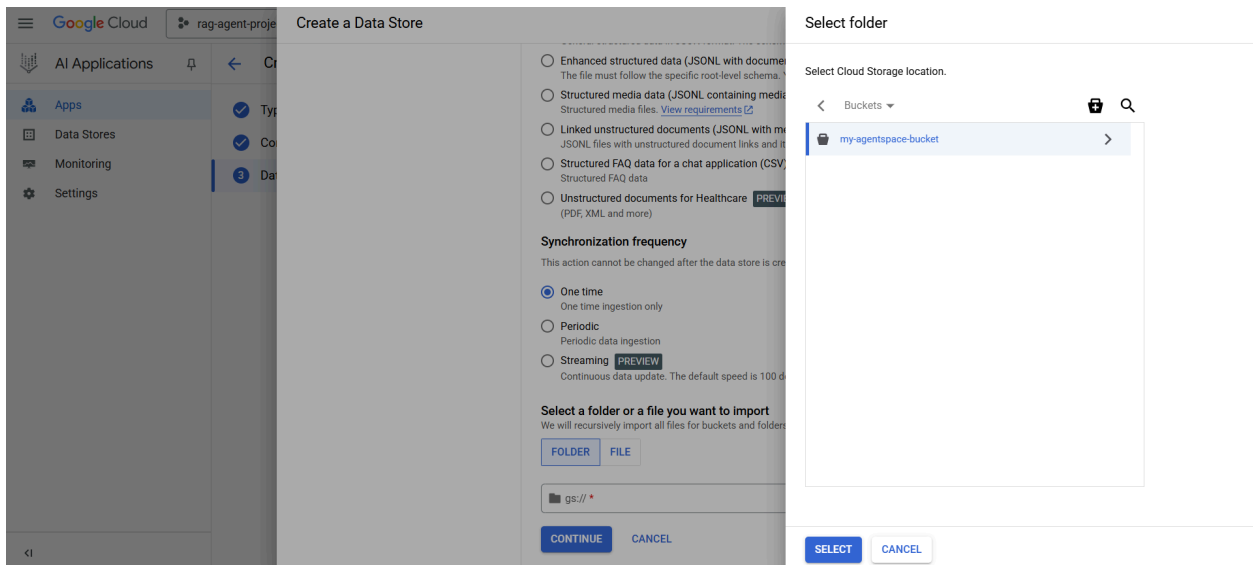
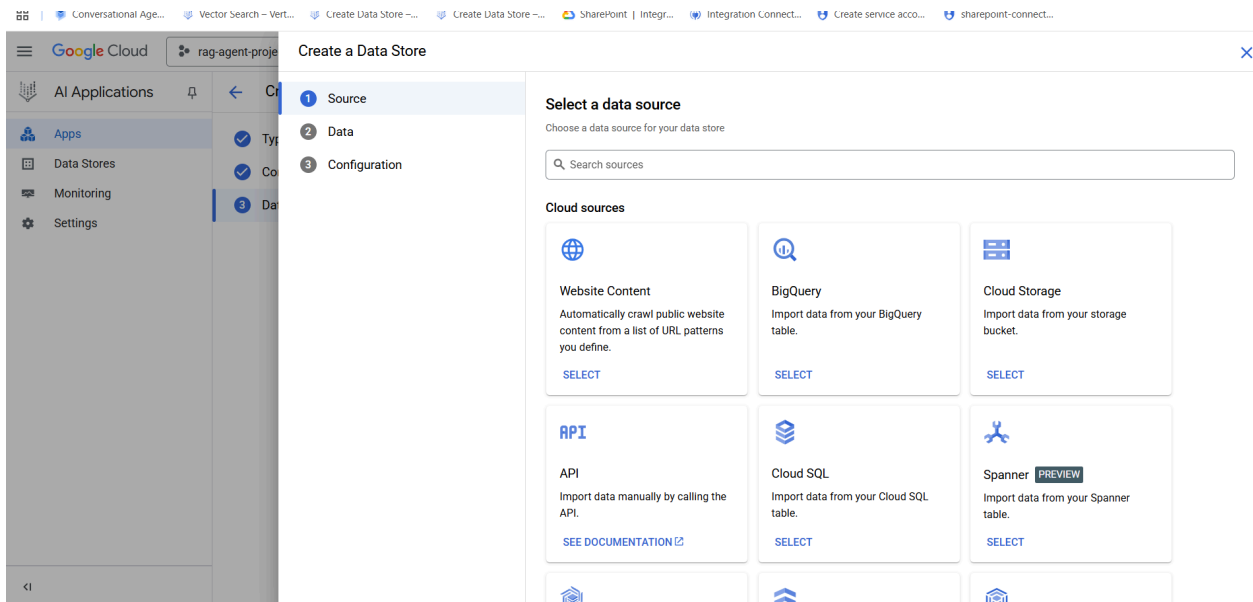
➔ Custom search (general)

This one is designed for:

- Unstructured content (like your PDFs)
- Cloud Storage connectors
- Generative AI mode with chunk + rerank
- And flexible chunking

The screenshot shows the 'Create App' interface in the Google Cloud Console for the 'rag-agent-project'. The left sidebar contains a menu with 'Apps', 'Data Stores', 'Monitoring', and 'Settings'. The main content area is titled 'Create App' and includes a 'Discovery Engine' tab. A warning message at the top states: 'Advanced generative answers. Advanced LLM features are not available for basic website search. You can change this setting at any time. After turning on Advanced LLM features, it can take up to 5 minutes for the features to become available. [Learn more about features and prices](#)'. Below this, the 'Your app name' section has a text input for 'App name' with the value 'rag-agent-app' and a note: 'ID: rag-agent-app_1751828087298. It cannot be changed later. [EDIT](#)'. The 'External name of your company or organization' section has a text input for 'Company name'. The 'Location of your app' section has a dropdown menu for 'Multi-region' with the selected option 'us (multiple regions in United States)' and a note: 'You can not change it later. For important information about multi-regions, see [Vertex AI Search locations](#)'. At the bottom, there are 'CONTINUE' and 'CANCEL' buttons.

lick CREATE DATA STORE (the blue button at the top).



✓ Default document parser

- leave as **Layout Parser** (good for PDFs with headings and sections)

✓ Layout parser settings

- *Optional:*

- Enable table annotation → check this if your PDFs have important tables you want the LLM to reason over
 - Enable image annotation → only check if you have images with text you want to extract (OCR).
 - If your PDFs are mostly text, you can leave both unchecked for now.
-

✓ Document chunking

- ✓ *Enable advanced chunking configuration* (already checked in your screenshot)
- Chunk size limit: 500 tokens → perfect
- Include ancestor headings in chunks → you can check this if you want the chunk to also include its parent headings (like a section name) for more context. Recommended for policy documents.

So I'd suggest:

- check "Include ancestor headings in chunks" ✓
(this helps provide extra context to the LLM)
-

✓ GENERATIVE AI OPTIONS (expand the accordion):

- Enable advanced generative answers
- Enable reranking

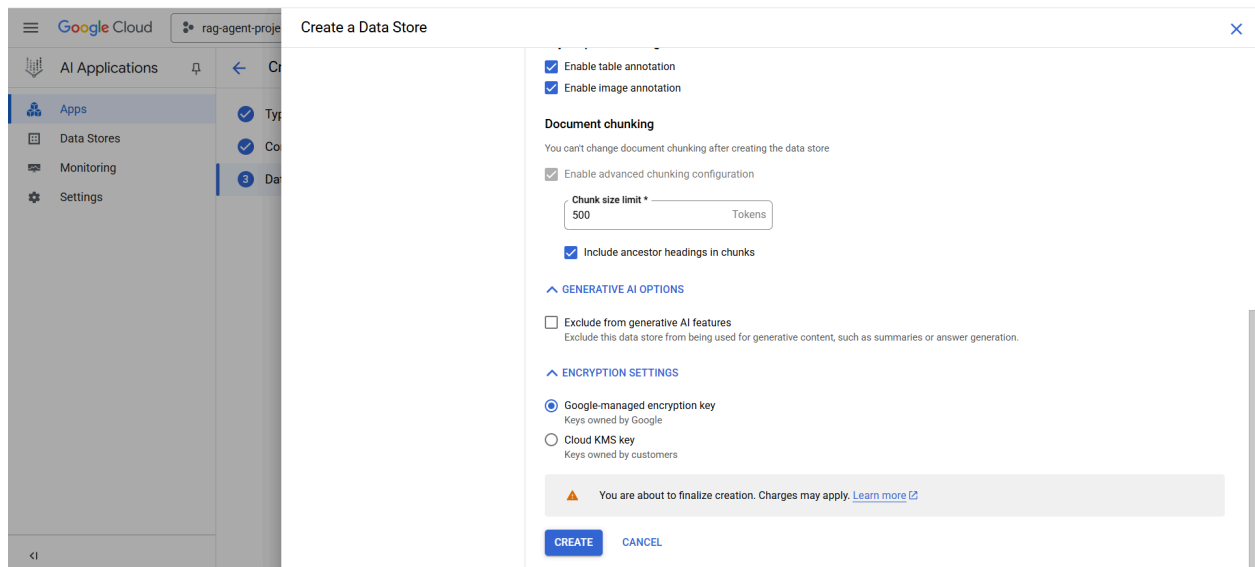
✓ ENCRYPTION SETTINGS

- you can leave default unless you have your own CMEK key

✓ Double-check:

- Default document parser: **Layout Parser**

- **Chunk size: 500**
- **Include ancestor headings: ☒ checked**
- **Table + image annotations: ☒ checked**
- **Advanced generative answers: active (leave “exclude” unchecked)**
- **Encryption: Google-managed**



- ✓ **Chunk size = 500 tokens (you set it)**
- ✓ **Overlap = 25% default (no box needed, built into the parser)**
- ✓ **Layout-aware = Layout Parser selected → ON**
- 👉 **This will:**
 - ✓ **finalize linking your app with the data store**
 - ✓ **kick off ingestion (you may see the status as **Ingesting** for a bit)**
 - ✓ **make your PDFs fully discoverable for questions**

Google Cloud

rag-agent-project

Discovery Engine

Search

3

?

:

A

AI Applications

Create App

LEARN

Apps

Data Stores

Monitoring

Settings

Advanced generative answers

Advanced LLM features are not available for basic website search. You can change this setting at any time. After turning on Advanced LLM features, it can take up to 5 minutes for the features to become available. [Learn more about features and prices](#)

Your app name

App name *

rag-agent-app

ID: rag-agent-app_1751828959200. It cannot be changed later. [EDIT](#)

External name of your company or organization

Company name *

rag-agent-lab

Providing your company name helps the model provide higher-quality responses

Location of your app

We recommend that you choose the global location, if you do not have compliance or regulatory reasons to locate your data in a particular multi-region. (EU and US regions are currently in preview)

Multi-region *

us (multiple regions in United States)

You can not change it later. For important information about multi-regions, see [Vertex AI Search locations](#)

CONTINUE

CANCEL

Google Cloud

rag-agent-project

Discovery Engine

Search

3

?

:

A

AI Applications

Create App

LEARN

Apps

Data Stores

Monitoring

Settings

Type

Configuration

Data

Data Stores

CREATE DATA STORE

Filter

Enter property name or value

?

☐

Name

Connected apps

Created ↓

ID

Location

☐

rag-agent-datastore

[rag-agent-app](#)

Jul 6, 2025

rag-agent-datastore_1751828318913

us

CREATE

CANCEL

Google Cloud

rag-agent-project

Discovery Engine

Search

3

?

:

A

AI Applications

Create App

LEARN

Apps

Data Stores

Monitoring

Settings

Type

Configuration

Data

Data Stores

CREATE DATA STORE

Filter

Enter property name or value

?

☒

Name

Connected apps

Created ↓

ID

Location

☒

rag-agent-datastore

[rag-agent-app](#)

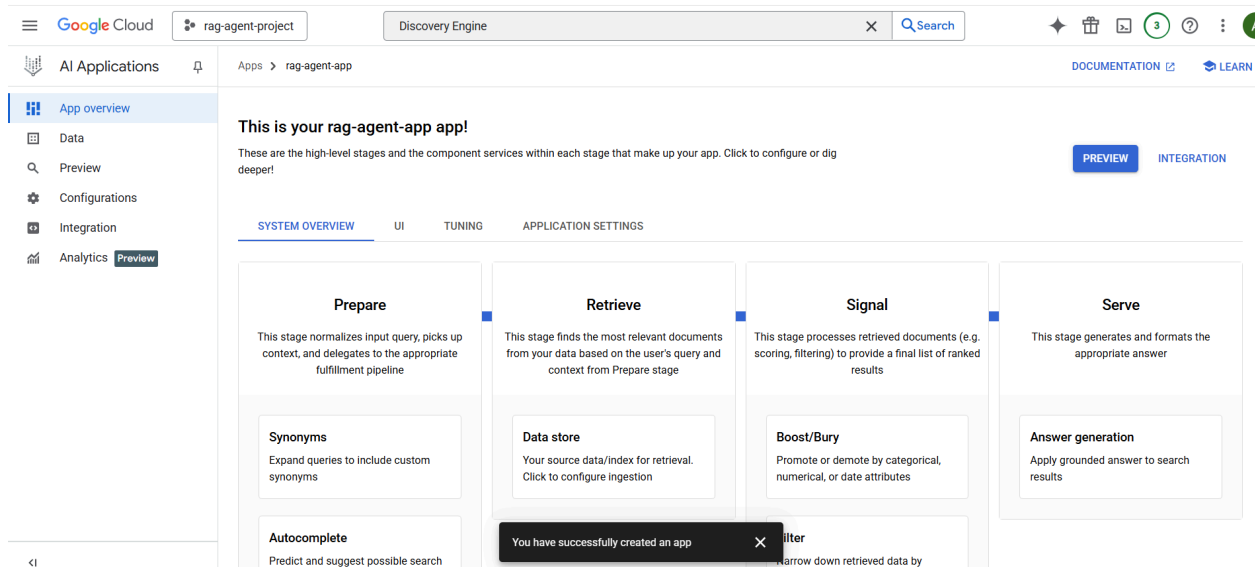
Jul 6, 2025

rag-agent-datastore_1751828318913

us

CREATE

CANCEL



Test your app

👉 Click the blue PREVIEW button (top right of your screenshot).

This will open the *Test UI* where you can send sample queries.

✅ The system will:

- run semantic search over your chunked documents
- rerank the top chunks
- send them to the LLM (Gemini or PaLM)
- generate an answer
- and highlight which chunks from your documents were used

Google Cloud

rag-agent-project

Discovery Engine

AI Applications

Apps > rag-agent-app > Data

App overview

Data

Preview

Configurations

Integration

Analytics Preview

rag-agent-datastore

Data store IDrag-agent-datastore_1751828318913

TypeUnstructured data

Serving stateEnabled

Regionus

LanguageN/A

Datastore size-

Number of documents-

Last document importJul 6, 2025, 3:04:05 PM

Exclude from generative AI featuresFalse

DOCUMENTS

EVENTS

ACTIVITY

PROCESSING CONFIG

PREVIEW

You have successfully created an app

Activity log details

StatusImport in progress

DetailsNo errors

Items succeeded0

Operation nameimport-documents-9819417407278449328

Last updatedSun Jul 06 2025 15:04:05 GMT-0400 (Eastern Daylight Time)

Google Cloud

rag-agent-project

Discovery Engine

AI Applications

Apps > rag-agent-app > Data

App overview

Data

Preview

Configurations

Integration

Analytics Preview

rag-agent-datastore

Data store IDrag-agent-datastore_1751828318913

TypeUnstructured data

Serving stateEnabled

Regionus

LanguageN/A

Datastore size-

Number of documents7

Last document importJul 6, 2025, 3:04:05 PM

Exclude from generative AI featuresFalse

DOCUMENTS

EVENTS

ACTIVITY

PROCESSING CONFIG

PREVIEW

+ IMPORT DATA

PURGE DATA

ID35b7677ed01401288b9e39931e9ea3f0

URIgs://m

You have successfully created an app

Activity log details

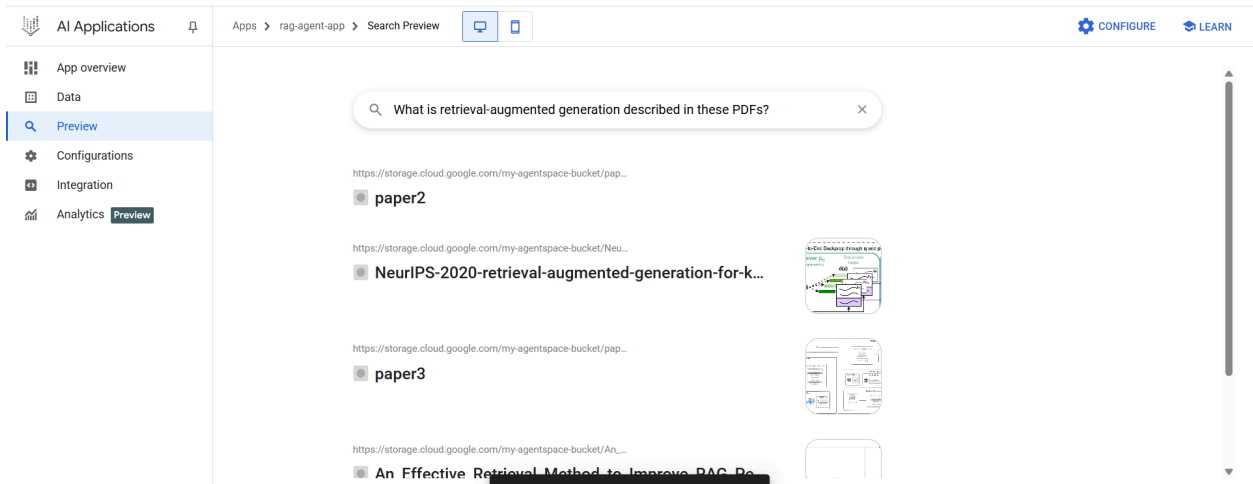
StatusImport completed

DetailsNo errors

Items succeeded7

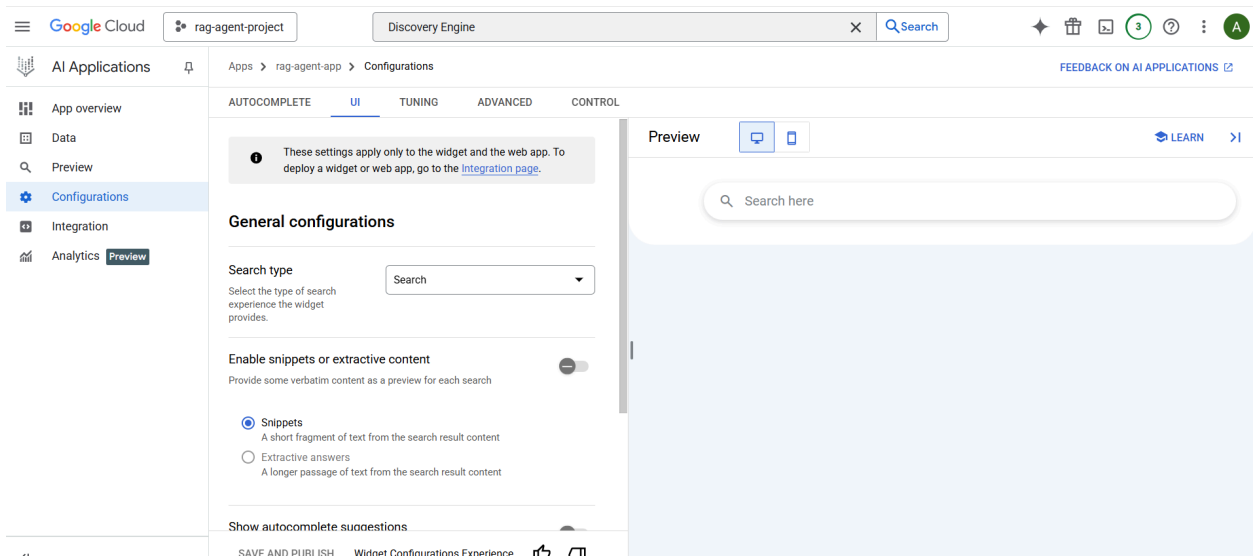
Operation nameimport-documents-9819417407278449328

Last updatedSun Jul 06 2025 15:17:36 GMT-0400 (Eastern Daylight Time)



Go to your app configuration

- Click on **App overview** → **Serve** → **Answer generation**
- Make sure “Grounded Answer Generation” is turned **on** (it usually is by default, but check).



To achieve your goal — specific answer with references + section names

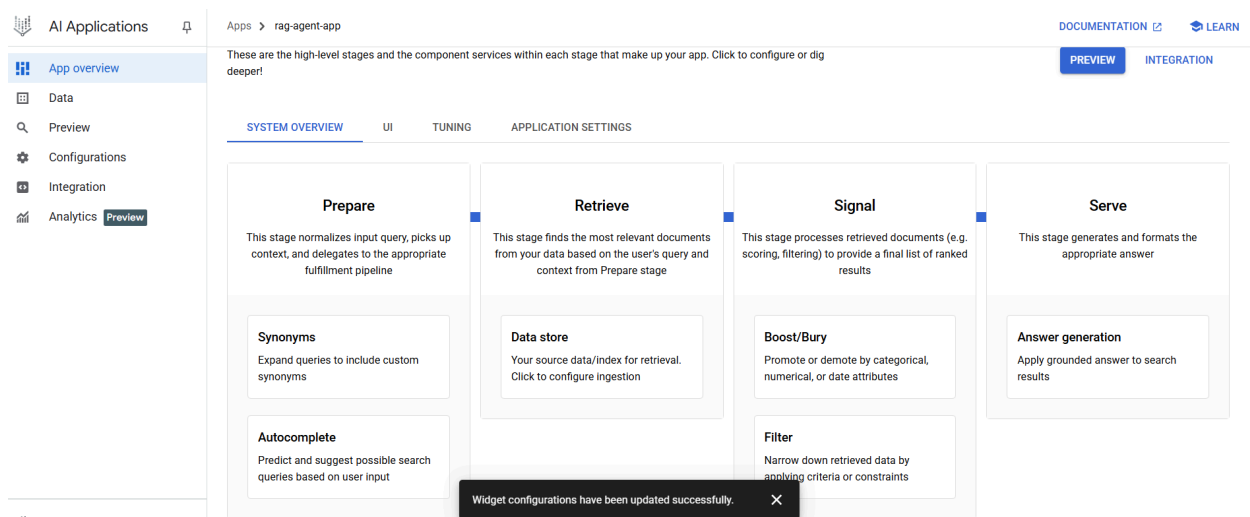
DO THIS INSTEAD:

- ✓ **Stay with Snippets** (required with advanced chunking)
- ✓ Then use the **Grounded Answer Generation** capability:

1. Go to **App Overview > Serve > Answer Generation**
 - enable *Grounded answers*
 - that lets the LLM combine chunks into a single natural-language answer
 - with references to each chunk (containing subtitle, heading, page etc)
2. Save and deploy

✓ **Result:**

- The system will use advanced chunking (as you set)
- then let the LLM *generate* a final answer with:
 - supporting chunks
 - and metadata (section names, page numbers, etc)



✓ **1. Save and publish**

Click the blue **Save and Publish** button at the bottom left, so these snippet settings take effect.

✓ 2. Go to the App Overview

- In the left menu, click **App overview**
- You'll see the four stages: *Prepare*, *Retrieve*, *Signal*, *Serve*
- In the **Serve** stage, click **Answer generation**

✓ 3. Enable Grounded Answer Generation

- Inside the *Serve* stage, there should be a toggle for **Answer Generation** (sometimes called “grounded answers” or “generative answers” in Google Discovery Engine)
- Enable it
- This feature allows the LLM to:
 - Take the chunks retrieved from your vector DB
 - Combine them
 - Generate a final answer
 - Include references to the chunk metadata (like heading, paragraph name, or page number)

👉 “Search with an answer”

Why?

- *Search with an answer* will generate a grounded answer above the search results, **and** it will still list the sources (chunks/paragraphs) with references.
- This mode allows you to get the best of both worlds:
 - A generative summary of the answer
 - While still surfacing the original passages with their metadata (section, page, heading)
 - And enables highlighted tokens in the result view

Recommendation: Select “Gemini 2.0 Flash 1”

because:

- ✓ tuned for Q&A and summarization
- ✓ fast (Flash)
- ✓ optimized for grounded answers
- ✓ supports multi-passage summarization with references

Next steps after selecting Gemini 2.0 Flash 1

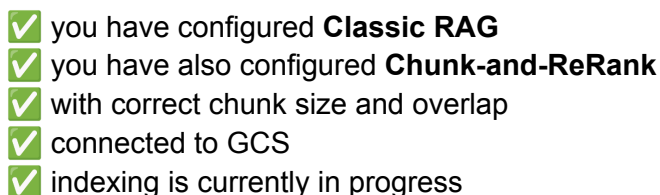
1. Save and Publish
2. Test in Preview
3. If needed, enrich your document metadata further (section, page, heading) during chunking
4. Verify highlighted references are showing up as expected

setting “English” might be safest, to keep all summaries uniform, unless you expect mixed languages.

👉 **My recommended setup for your case (testing + highlighting tokens + paragraph metadata):**

- **Enable related questions: ON** ✓
- **Ignore no answer summary for query: OFF** ☐ (so you see when no relevant answer exists)
- **Ignore Adversarial Query: OFF** ☐ (OK while testing)

When you go to production, you might revisit “Ignore Adversarial Query” and turn it ON.



yes i want: ☒ answers pointing to specific paragraphs/subtitles ☒ highlighted tokens ☒ and paragraph metadata (like "section name", "page", "heading")

- When ingesting documents, *chunk* them with a *custom chunker* that adds metadata for each chunk, like:
 - `page_number`
 - `section_heading`
 - `paragraph_index`

- `original_text`
- Store these fields as metadata in the vector database (Discovery's vector store).

✓ Then, in your retrieval or answer generation step, you will:

- pull back the top relevant chunks
- include their metadata in the LLM prompt (so it can cite sections)
- and optionally **highlight** the best matching tokens in the UI, if you build a frontend over the API.

2 How to inject metadata during chunking

If you are chunking a PDF with `Layout Parser` in Discovery, you cannot easily add advanced metadata in the console — so you'll want to do **custom ingestion** with something like a Python script.

Sample chunk + metadata:

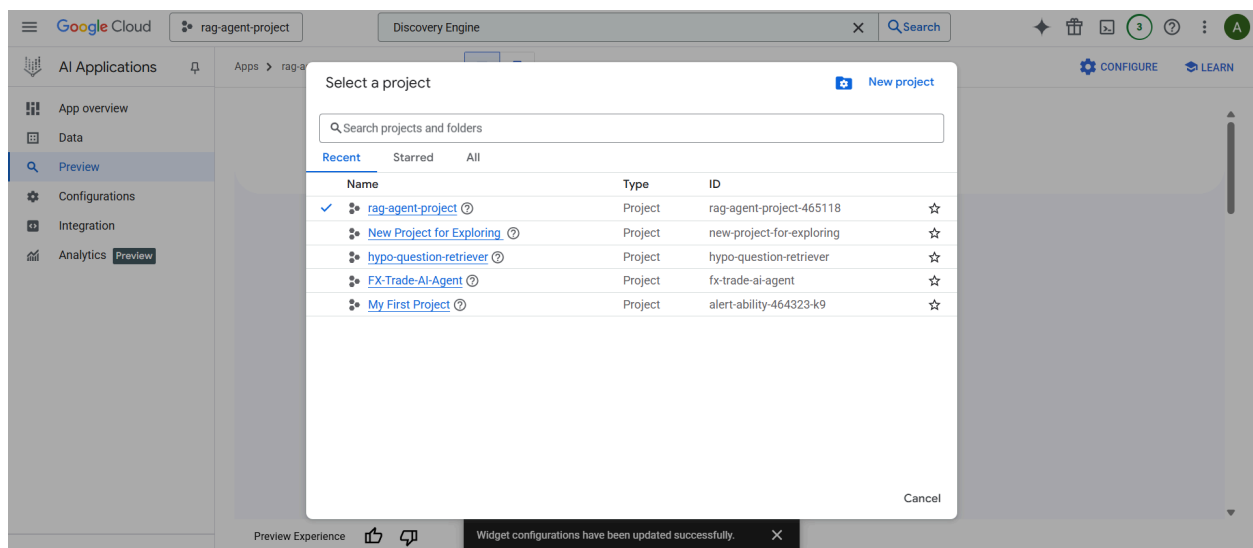
```
python
CopyEdit
{
  "id": "doc_123_page5_para2",
  "embedding": [0.45, 0.78, ...],
  "text": "Our employee vacation policy covers full-time
employees...",
  "metadata": {
    "source": "wells_handbook.pdf",
    "page": 5,
    "section": "Time Off and Leave",
    "heading": "Vacation Policy",
    "paragraph_index": 2
  }
}
```


You can do this chunking with:

- ✓ PyPDF2 / pdfminer to split PDF
- ✓ add headings from layout parsing
- ✓ embed with Gemini Embeddings API
- ✓ store with Discovery's **dataConnector** API (or even store in a hybrid vector db like Pinecone if you prefer, and connect Discovery on top)

Let me be very clear: **most** of these advanced metadata chunking steps happen *outside* the Google Cloud Console, because the console itself does not do “paragraph-level metadata enrichment” automatically. But I will explain **where** the console fits in each step so you have the full picture.

- ✓ parse a PDF
- ✓ extract headings
- ✓ chunk with metadata
- ✓ embed with Gemini
- ✓ upload to Discovery



Google Cloud

rag-agent-project

Search (/) for resources, docs, products, and more

Q Search

Cloud Storage

Bucket details

Go to path Refresh

Overview

Buckets

Monitoring

Settings

Storage Intelligence

Insights datasets

Configuration

my-agentspace-bucket

Location: us (multiple regions in United States)

Storage class: Standard

Public access: Not public

Protection: Soft Delete

Objects

Configuration

Permissions

Protection

Lifecycle

Observability

New

Inventory Reports

Operations

Folder browser

my-agentspace-bucket

Create folder Upload Transfer data Other services Learn

Filter by name prefix only Filter Filter objects and folders Show Live objects only

	Name	Size	Type	Created	Storage class	
<input type="checkbox"/>	An_Effective_Retrieval_Method_to...	270.6 KB	application/pdf	Jul 6, 2025, 2:48:23 PM	Standard	Download More actions
<input type="checkbox"/>	NeurIPS-2020-retrieval-augmented...	946.1 KB	application/pdf	Jul 6, 2025, 2:48:25 PM	Standard	Download More actions
<input type="checkbox"/>	borgeaud22a.pdf	986.9 KB	application/pdf	Jul 6, 2025, 2:48:26 PM	Standard	Download More actions
<input type="checkbox"/>	paper1.pdf	6.3 MB	application/pdf	Jul 6, 2025, 2:48:40 PM	Standard	Download More actions
<input type="checkbox"/>	paper2.pdf	199.2 KB	application/pdf	Jul 6, 2025, 2:48:30 PM	Standard	Download More actions
<input type="checkbox"/>	paper3.pdf	11 MB	application/pdf	Jul 6, 2025, 2:48:43 PM	Standard	Download More actions
<input type="checkbox"/>			application/pdf	Jul 6, 2025, 2:48:50 PM	Standard	Download More actions

Now viewing project "rag-agent-project" in organization "No organization"

Google Cloud

rag-agent-project

Vertex AI Model Garden

Q Search

Model Garden

Explore Generative AI View my endpoints & models Deploy from Hugging Face View release notes

Tasks

Text embeddings 1

Model Collections

Google models 1

Providers

Google 1

AI2 1

Features

Search models

textembedding-gecko

Search results for "textembedding-gecko"

Sort by: Trending Newest Last updated

Name	Description	Type	Task	Modality	Resource ID
Embeddings for Text	Converts text data into vector representations for semantic search,...	Foundation	Embedding	Language	text-embedd...

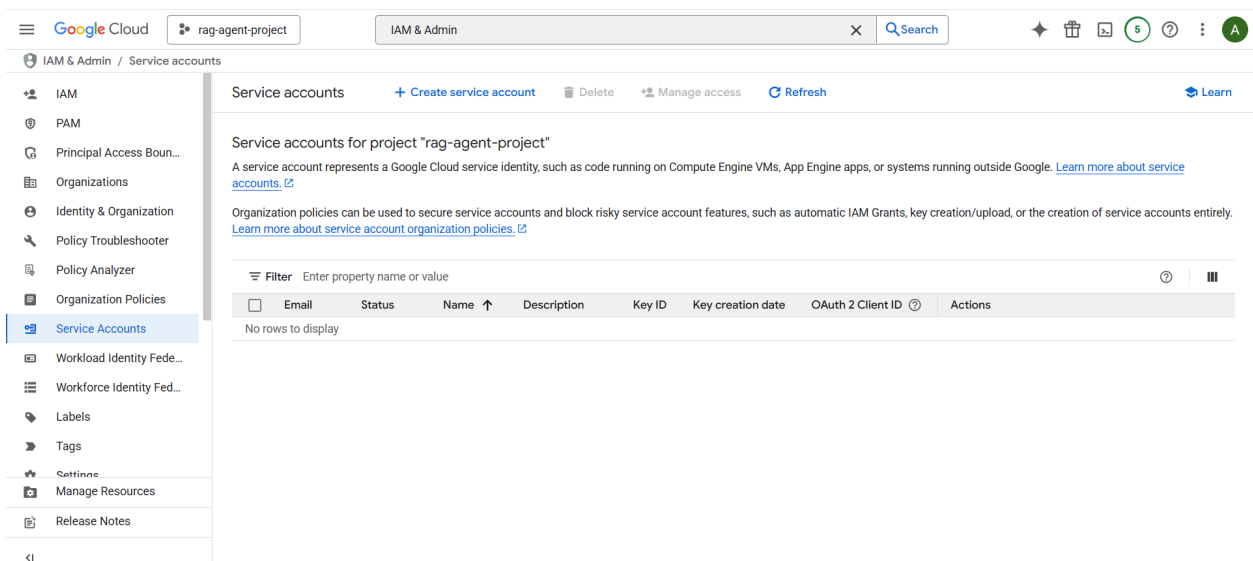
[https://cloud.google.com/python/docs/reference/aipatform/1.28.0/vertexai.language_models.TextEmbeddingModel#:~:text=Dismiss%20View.:%20str\)%20%2D%3E%20vertexai.](https://cloud.google.com/python/docs/reference/aipatform/1.28.0/vertexai.language_models.TextEmbeddingModel#:~:text=Dismiss%20View.:%20str)%20%2D%3E%20vertexai.)

You obtain the `/path/to/your/service_account.json` file when you create and download a service account key from your Google Cloud project.

Here's a step-by-step guide:

1. **Go to the Google Cloud Console:**
 - Open your web browser and navigate to <https://console.cloud.google.com/>.
 - Make sure you select the correct Google Cloud project where you want to use Vertex AI.
2. **Navigate to IAM & Admin -> Service Accounts:**
 - In the Google Cloud Console, use the navigation menu (usually on the left) and go to "IAM & Admin" > "Service Accounts."
3. **Create a New Service Account (if you don't have one):**
 - Click on "+ CREATE SERVICE ACCOUNT" at the top of the page.
 - **Service account name:** Give it a descriptive name (e.g., `vertex-ai-user`, `my-app-service-account`).
 - **Service account ID:** This will be automatically generated based on the name.
 - **Service account description:** (Optional) Add a brief description.
 - Click "CREATE AND CONTINUE."
4. **Grant Permissions/Roles:**
 - This is the crucial step for giving your service account the necessary access to Vertex AI.
 - In the "Grant this service account access to project" section, click on the "Select a role" dropdown.
 - Search for and add the following roles (at a minimum):
 - `Vertex AI User`
 - `Service Usage Consumer` (this allows the service account to use APIs enabled in your project)
 - Depending on your specific needs, you might also consider:
 - `Vertex AI Developer` (broader permissions for development)
 - `Storage Object Viewer` or `Storage Object Creator` (if your application needs to read/write from Cloud Storage buckets)
 - **Principle of Least Privilege:** Always grant only the necessary permissions. Avoid giving roles like "Owner" or "Editor" unless absolutely required, especially in production environments.
 - Click "CONTINUE."
5. **Grant users access to this service account (Optional):**
 - You can skip this step unless you need to grant other users or service accounts the ability to *impersonate* this new service account.
 - Click "DONE."
6. **Create and Download the JSON Key:**
 - Now that the service account is created, you need to generate a key for it.

- On the Service Accounts page, find the service account you just created.
 - Click on the three vertical dots (Actions menu) under the "Actions" column for your service account.
 - Select "Manage keys."
 - Click on "ADD KEY" > "Create new key."
 - Select "JSON" as the key type.
 - Click "CREATE."
7. Your browser will automatically download a JSON file. This file contains the private key for your service account and is what you'll use for authentication.
8. **Store the JSON Key Securely:**
- **This file is highly sensitive.** Treat it like a password. Anyone who has this file can authenticate as your service account and access resources it has permissions for.
 - **Do NOT** commit it to version control (Git, etc.).
 - Store it in a secure location on your machine or server where your code will run.
9. **Set the Environment Variable:**
- Once you have the JSON file, replace `/path/to/your/service_account.json` in your code with the actual path to the downloaded JSON file on your system.
10. For example, if you downloaded `m`



Google Cloud

rag-agent-project

IAM & Admin

Search

IAM & Admin / Service accounts / Create service account

IAM

PAM

Principal Access Boun...

Organizations

Identity & Organization

Policy Troubleshooter

Policy Analyzer

Organization Policies

Service Accounts

Workload Identity Fede...

Workforce Identity Fed...

Labels

Tags

Settings

Manage Resources

Release Notes

Create service account

1 Create service account

Service account name

discovery-admin

Display name for this service account

Service account ID *

discovery-admin

X

↺

Email address:

discovery-admin@rag-agent-project-465118.iam.gserviceaccount.com

📧

Service account description

Discovery and VertexAI access for RAG application

Describe what this service account will do

Create and continue

2 Permissions (optional)

3 Principals with access (optional)

Done

Cancel

Google Cloud

rag-agent-project

IAM & Admin

Search

IAM & Admin / Service accounts / Service account: 107930932126774237119 / Permissions

IAM

PAM

Principal Access Boun...

Organizations

Identity & Organization

Policy Troubleshooter

Policy Analyzer

Organization Policies

Service Accounts

Workload Identity Fede...

Workforce Identity Fed...

Labels

Tags

Settings

Manage Resources

Release Notes

discovery-admin

Details

Permissions

Keys

Metrics

Logs

Manage service account permissions

You can edit roles assigned to a service account on resources in t manage access to resources on other projects in your organizatio page.

Manage access

View service account permissions

You can use Policy Analyzer to view which resources this service access to.

Run Policy Analyzer

Principals with access to this service acco

1

Principals can be granted access to impersonate servi

Edit access to "rag-agent-project"

Principal

discovery-admin@rag-agent-project-465118.iam.gserviceaccount.com

rag-agent-project

Assign roles

Roles are composed of sets of permissions and determine what the principal can do with this resource. [Learn more](#)

Role

Discovery Engine Admin

IAM condition (optional)

+ Add IAM condition

Grants full access to all discoveryengine resources.

Role

Vertex AI User

IAM condition (optional)

+ Add IAM condition

Grants access to use all resource in Vertex AI

Role

Storage Object Viewer

IAM condition (optional)

+ Add IAM condition

Grants access to view objects and their metadata, excluding ACLs. Can also list the objects in a bucket.

+ Add another role

Save

Test changes

Cancel

Summary of changes

Roles removed

n/a

Roles added

Storage Object Viewer

Vertex AI User

Test changes

Google Cloud

rag-agent-project

IAM & Admin

Search

IAM & Admin / Service accounts

Service accounts

+ Create service account

Delete

Manage access

Refresh

Learn

Service accounts for project "rag-agent-project"

A service account represents a Google Cloud service identity, such as code running on Compute Engine VMs, App Engine apps, or systems running outside Google. [Learn more about service accounts](#).

Organization policies can be used to secure service accounts and block risky service account features, such as automatic IAM Grants, key creation/upload, or the creation of service accounts entirely. [Learn more about service account organization policies](#).

Filter

Enter property name or value

🔍

⌵

<input type="checkbox"/>	Email	Status	Name ↑	Description	Key ID	Key creation date	OAuth 2 Client ID	Actions
<input type="checkbox"/>	discovery-admin@rag-agent-project-465118.iam.gserviceaccount.com	Enabled	discovery-admin	Discovery and VertexAI access for RAG application	No keys		107930932126774237119	⋮

Policy updated

Perfect — you are **almost there**, but notice it says:

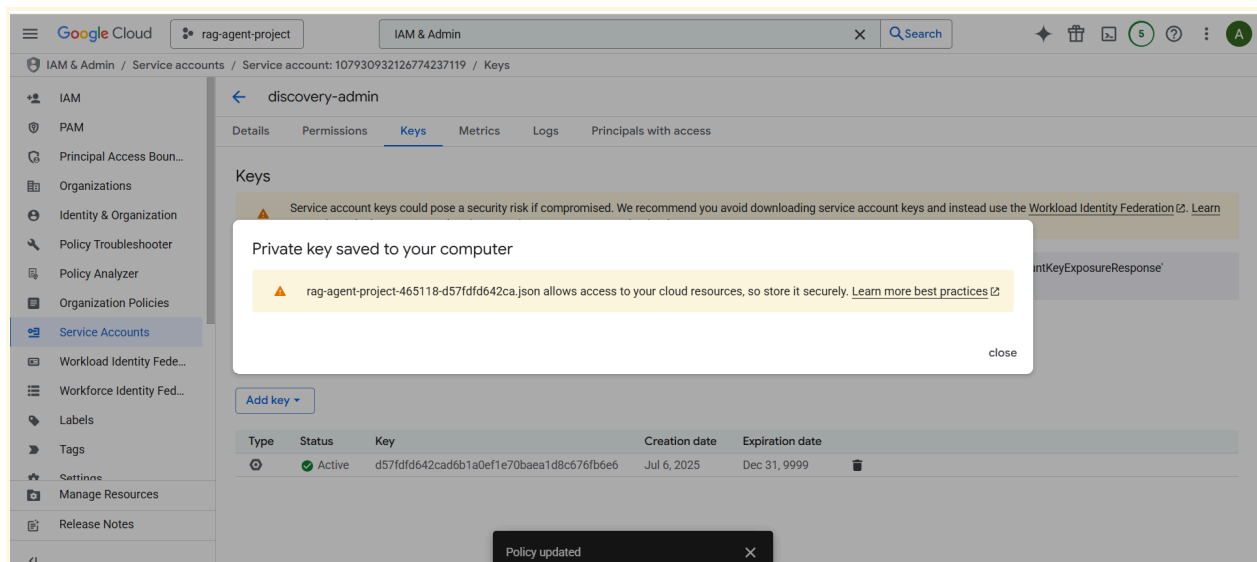
Key ID → *No keys*

✅ That means you still need to generate a key file so you can authenticate from Colab or local.

Do this next (exact steps):

- ❶ Click the **three dots** on the right of `discovery-admin@...`
- ❷ Select **Manage keys**
- ❸ Click **Add Key** → **Create new key** → **JSON**
- ❹ Download the JSON key file (for example: `discovery-admin-key.json`)
- ❺ Upload that JSON file to Colab
- ❻ In Colab, set:

`rag-agent-project-465118-d57dfd642ca.json`



🚀 Awesome — let's go step by step to **push your chunks to Discovery**.

✅ **High-level steps:**

1. You already have `all_chunks` with text + metadata + embeddings
2. Now we'll use the Discovery Engine Data API to push these chunks
3. Then you'll verify the chunks are indexed and ready for RAG

